

000  
001  
002  
003  
004  
005  
006  
007  
008  
009  
010  
011  
012  
013  
014  
015  
016  
017  
018  
019  
020  
021  
022  
023  
024  
025  
026  
027  
028  
029  
030

## Abstract

Style transfer aims at creating new images that preserve from a sample image the subjective characteristics known as style. The transfer can be considered as successful if the new image clearly preserves the composition of the original image while producing the same visual impression as the image from which we imitate the style. With the development of virtual and augmented reality, there is a growing demand for high quality synthesized images that produce either realistic or graphical impressions on the viewer. By providing realistic or style-specific visual results on a finer scale to basically drafted synthesized images, style transfer can significantly enrich the user's visual experience. Another interest is purely artistic, as it allows to create new images that mimic various painting styles, therefore producing interesting aesthetic results. One approach of style transfer is directly inspired from texture transfer and aims at mapping patches from the style image to the content image in order to preserve small scale details while trying to produce comparable large scale items. An inherent problem in this case is therefore the ability to distinguish between object details and style attributes that are of similar size. In this paper, we further explore the application of adaptive size patch quilting to style transfer to preserve fine-grained details of the original image that are below the maximal scale of the texture pattern. To do this, we use a quad-tree structure that takes into account both the local homogeneity of the content image and the local ability to be faithfully mapped to the style image. We try to improve upon the original method to allow for more fine-tuning of the splitting. Allowing for more adaptation of the quad-tree splitting for the transfer task given a specific style-content pair of images improved the final visual result. We also propose to operate the regularization of the patch selection at the same time as the splitting and quilting occurs, using the same rules for the splitting as for the patch selection. This allows us to have a splitting that takes into account all the constraints of the problem.

## 1 Introduction

In a style transfer task, two uncorrelated images are used to produce a new one. A 'content' image contains a layout of recognizable items, which form the composition of the scene. A second 'style' image serves as an example of the general visual impression that we want the final image to convey. One of the difficulties is that the style image usually depicts recognizable items beyond the specific style. The extraction of style attributes from the style image must therefore be conducted carefully in order to capture the style's characteristics independently from composition of the style image.

Although style refers to an immediate visual perception, it is non-trivial to formalize. A first approximation is to consider style as the lower-scale details of the image that are independent of the content, and therefore found repeatedly in places that correspond to different items of the image. For instance, it can be the aspect of the brush-stroke or a pencil, or larger-scale patterns that are consistently repeated by the painter across the painting.

In this context, texture-transfer can be seen as a specific case of style-transfer where the style image does not contain any recognizable items and is therefore restricted to texture details. As texture transfer is a problem that has been extensively studied, style transfer methods can directly benefit from the methods that were developed for this purpose.

For texture-transfer, various approaches have produced visually satisfying results.

Methods that tried to extensively capture the statistics of the underlying stochastic process have been rapidly rejected because of the important computational effort involved by such methods. The first visually interesting results in texture transfer derive from techniques where the new image is grown one pixel at the time.

Wei and Levoy [8] considered the entire possible span of pixels' neighborhoods from the texture image to grow textures one pixel at the time, selecting the most likely color values given the already synthesized neighborhood, which is named the causal neighborhood. They also introduced a

054 multi-scale pyramidal synthesis process which allowed high-frequency details to be coherent with  
055 low-frequency ones. By synthesizing images forming a pyramid from a coarse scale to finer ones,  
056 considering the previously synthesized at various coarse scales, they achieved state of the art results  
057 for texture synthesis. They therefore created an algorithm that allowed to create larger samples of  
058 a texture from an original small sample, thus opening the way to texture transfer techniques where  
059 textures were mapped pixel by pixel from a small samples to larger example images.

060 The described method has the drawback of being computationally expensive and sometimes gets  
061 stuck in local minima, synthesizing always the same sequence of pixels and therefore failing to re-  
062 construct complex patterns by producing 'garbage'. The second limitation is solved by Ashikhmin  
063 [1] which also synthesizes the pixels one by one in scanline order. The new pixel is chosen from  
064 a small list of candidates by minimizing neighborhood distances. The candidates are selected by  
065 looking at the pixels from the texture image that were mapped to the causal neighborhood of the  
066 currently synthesized pixel. Among those neighbors, the chosen pixel is the one that has the same  
067 relative positioning to the causal pixel and for which the neighborhood distance is minimal. Pre-  
068 serving the relative positioning of the candidate to the neighbor pixel in the final image is effectively  
069 equivalent to growing patches of irregular shapes. This method faithfully transfers the texture in  
070 the specific case of textures composed of small irregular shapes of various small sizes, but produces  
071 sharp edges between two synthesized patches which makes this method inadequate for general tex-  
072 ture. Ashikhmin [2] also introduced a variant that performs texture transfer onto simple flat color  
073 drawings, getting closer to our more complex objective of transferring styles. This method jointly  
074 minimizes the distance of the causal neighborhood to the texture sample image and the comple-  
075 mentary part of the neighborhood to the color image. The algorithm therefore overwrites the color  
076 image with pixels from the texture sample by preserving both a local coherency between already  
077 synthesized pixels and a consistency with the color image neighborhood.

077 A somewhat different approach is taken in [3], as it introduces the idea of mapping square image  
078 patches sampled from the texture image to the content image by minimizing the intensity distances.  
079 This approach therefore does not rely on individual pixel synthesis and can be computationally more  
080 efficient. In this case, the quality of the transfer depends strongly on the chosen size of the patch,  
081 which is taken as a rule of thumb as the order of magnitude of the largest visible patterns of the  
082 texture. Using this patch size focuses on preserving the full complexity of the pattern, but often  
083 at the expense of the finer details of the original image. As patches are used, borders have to be  
084 blended in order to avoid introducing a visually striking square grid on the final image. Efros and  
085 Freeman [3] therefore performs the image quilting with an overlap and operate a boundary cut in  
086 order to attenuate the square's borders. The borders are made uneven in a way that maximizes visual  
087 coherency, operating the cut where the pixels are most similar in the overlapping region. Visually,  
088 this method often outperforms methods that focus on individual pixels at a lesser computational  
089 price. In this article, although style transfer is not explicitly mentioned, the method could be directly  
090 applied to this task. In this case, the same shortcomings as the ones for texture transfer would apply  
091 because of the fixed choice of the patch size.

092 A new method presented by Frigo et al [4] explicitly focuses on style transfer and tries to alleviate  
093 the issue of a fixed-size patch by introducing an adaptative quilting technique. Instead of having a  
094 fixed path size, a quadtree structure is constructed before the quilting is performed. Starting from a  
095 square image, successive splits are made, each split dividing a square block into four equal smaller  
096 squares. The quadtree construction takes into account two factors in order to decide whether a finer  
097 split should occur. The first parameter considered is the local homogeneity of the content image,  
098 which is computed by the mean of the pixel variance of the current square block of the content  
099 image. Favoring local homogeneity encourages splitting if the current square is irregular, and is  
100 therefore likely to contain relevant content details. The second parameter is the local proximity to  
101 the style image, which is accounted for by the minimal squared pixel distance to same-size patches  
102 of the style image. This way, splitting occurs if no satisfying match for the current patch is found in  
103 the style image. As different patch sizes are considered, a normalized patch distance  $d$  is considered  
104 by dividing the sum of square of pixel distances over the patch by the number of pixels in the given  
105 block. The quadtree splitting also takes into account a minimal and a maximal patch size,  $\gamma_{min}$   
106 and  $\gamma_{max}$  which enables some balancing between the maximal size of exact copying of the original  
107 image, and the computational intensiveness. During the splitting, a set of optimal candidates from  
108 the style image are chosen from the style image using a constrained k-nn to provide a diversity  
109 among the selected patches. At the final stage a label is assigned to each block of the final quadtree

108 by taking into account both the closeness to the content block and the visual continuity between  
 109 style patches in order to try to provide the best visual result.  
 110

111 Building upon this work, we tried to improve the obtained result by addressing two aspects of the  
 112 quadtree quilting method. A first step is to separate the contribution of the local inhomogeneity and  
 113 the contribution of the distance to the style patches. This separation would allow us to dissociate  
 114 and finetune for a given image the persistence of the details and the closeness of the mapping to the  
 115 style image. We also decided to replace the constrained k-nn, choosing a single label at split and  
 116 quilting time and trying to compensate for the potential loss of continuity by operating a boundary  
 117 cut as presented by Efros and Freeman [3]. We therefore intend to produce an efficient merging  
 118 of the patches in cases where the two patches that are mapped as neighbors are not as close as in  
 119 the examples provided by the original quadtree quilting method, and the bilinear blending therefore  
 120 proves to be unsatisfying. On the other hand, to account for the effort to limit repetition of similar  
 121 patches close by, we will take into account the distance between blocks that have been mapped to  
 122 the same patch.

123 We will then compare our results to recent state of the art style transfer results that were introduced  
 124 by Gatys and al [5] using convolutional neural networks. This method minimizes a weighted loss  
 125 of composed of two terms, one that is associated to the style of an image, and the other that is  
 126 associated to the content of another image. The goal of the method is then to find an image that  
 127 minimizes this loss, effectively producing an image that, to some extent, imitates the style of the  
 128 first image while preserving the content of the second one. The style loss is expressed as a weighted  
 129 sum over several layers. It is taken to be the sum of square distances of the correlations between  
 130 activations for different filters across all possible filter positions. The content loss is simply the  
 131 sum of squared differences for a given convolutional layer between the activations for two different  
 132 images. By jointly minimizing these two losses on various layers, Gatys and all manage to transfer  
 133 styles from one image to another. The choice of the layers on which they choose to apply the content  
 134 loss determines to some extent the minimal scale at which the content will be preserved, with the  
 135 choice of earlier layers providing finer details.  
 136

## 137 2 Original adaptative quilting

138 The decision for a given square patch to be split can be summed up according to [4] by the following  
 139 condition :

$$140 \quad C(p_{x_i}, p_{y_i}) = ((\sigma_i + d(p_{x_i}, p_{y_i})) > \omega \text{ and } \tau_i > \gamma_{min}) \text{ or } (\tau_i < \gamma_{max})$$

142 Where  $d(p_{x_i}, p_{y_i})$  is the distance of the  $i^{th}$  patch of the content image to the closest patch  $p_{y_i}$  from  
 143 the style image.  $\tau_i$  being the current patch scale, the size in pixels of the side of the square patches.  
 144

145  $p_{y_i}$  is therefore the patch of the style image that minimizes the following distance :

$$146 \quad d(p_{x_i}, p_y) = \frac{\|p_{x_i} - p_y\|}{\tau_i^2}$$

149 Where the square of the chosen norm that is the sum of square differences for each color channels.  
 150

151 The variance term is the standard deviation of pixel illumination.

$$152 \quad \sigma_i = \sqrt{Var(p_{x_i})}$$

$$155 \quad \sigma_i = \sqrt{\sum_{x \in p_{x_i}} (I_x - I_m)^2}$$

158 where  $I_m$  is the mean pixel intensity,  $I_x$  is the illumination at pixel x, x belonging to the patch  $p_{x_i}$   
 159

160 Three parameters are left for the user of the algorithm to tune.  $\omega$ , which triggers a split if the sum  
 161 of the contributions of the local variance and of the distance to the style image are too large, and the  
 minimum and maximum patch sizes  $\gamma_{min}$  and  $\gamma_{max}$ .

162 For the leaf blocks of the quadtree (blocks which have failed the condition C and have therefore  
 163 reached their final size), k-nearest neighbors from the style image are computed. The k-nn algo-  
 164 rithm is constrained such as to limit the number of direct neighbors from the style labels among the  
 165 candidates. The constraint enforces a minimal distance between two candidate patches.

166 Once all candidate patches are computed for each of the leaf blocks, the final stitching occurs, trying  
 167 to optimize a probability density that considers for each patch a data fidelity term and a pairwise  
 168 compatibility term.

169 The term we want to maximize over the candidate label assignments is

$$172 \quad P(L) = \frac{1}{Z} \prod_i \Phi(l_i) \prod_{(i,j \in N)} \Psi(l_i, l_j)$$

175 Where i describes the leaf blocks' indexes, and N is the ensemble of direct neighbors in the quadtree  
 176 knowing that two blocks are considered as direct neighbors if one of their edges is at least partly  
 177 shared.

178  $\Phi$  is the data fidelity term, which takes into account the distance between the  $i^{th}$  block and the  
 179 candidate label.

$$182 \quad \Phi(l_i) = \exp(-\lambda_d * d(p_{x_i}^{imcontent}, p^{imstyle}l_i))$$

184  $\Psi$  is the pairwise compatibility term,

$$186 \quad \Psi(l_i, l_j) = \exp(-\lambda_s * d(\hat{p}_{l_i}^{imstyle}, \hat{p}_{l_j}^{imstyle}) + \lambda_r * |l_i - l_j|^2))$$

188  $\hat{p}_{l_i}^{imstyle}$  and  $\hat{p}_{l_j}^{imstyle}$  are the overlapping parts of two patches that have been grown in order to pro-  
 189 duce some overlap around the shared edge. This term aims at encouraging smoothness at the edges  
 190 as it enforces similarity between the patches around their shared limits. The second term penalizes  
 191 label repetition, as  $\Psi$  decreases if the two labels are very different.  $|l_i - l_j|^2$  is the square of the  
 192 distance between the two patches from the style image. We can notice that by penalizing label rep-  
 193 etition with this distance, there is also a penalization for close neighbors, therefore it is less likely  
 194 that direct neighbors in the style image will be mapped to direct neighbors in the final image. This  
 195 might be an unwanted side effect, as keeping neighbors together would increase perceived visual  
 196 coherence.

197 As patches are stitched together in this method, the issue of apparent edges has to be addressed. In  
 198 the original paper, a simple bilinear blending is applied, producing satisfying results in the presented  
 199 final images. Although bilinear blending is known to have limited smoothing capacities for images,  
 200 as the regularization emphasizes similarity over overlapping areas, the final bilinear blending is often  
 201 enough to suppress visible edges.

### 203 3 Improvements of quadtree style transfer

#### 205 3.1 Disentangling of variance and style proximity

207 To improve upon the method presented by Frigo and al [4], we started by reimplementing the  
 208 quadtree splitting algorithm as described in the original article in order to apply it to a new selection  
 209 of style and content images.

210 Our first observation was that using a quadtree structure as described by Frigo and al forces the  
 211 original content image and the final image to be square, and even more restrictively, the finale size  
 212 S of the image must be the product of  $\gamma_{min}$  and  $2^n$  where  $n \in \mathbb{R}$ . We therefore have  $S = \gamma_{min} * 2^n$ ,  
 213 and n is the maximum depth of splits that will occur during the quadtree construction.

215 We will therefore work exclusively with square images that respect this size constraint, cropping  
 and resizing the original content images when needed.

216 Starting with the square image, we recursively split each square into four squares until we reach  
217 the maximum size for a split  $\gamma_{max}$ . Once we've reached this block size, splits only occur if the  
218 condition for a split is satisfied, and the minimal split size  $\gamma_{min}$  hasn't yet been reached.  
219

220 To each square that has reached its final size in the content image, we want to assign a label that  
221 corresponds to a patch of the same size in the style image.

222 Implementing the original algorithm revealed an intrinsic limitation of the condition that is evaluated  
223 to determine whether a split to occur. By having a single parameter,  $\omega$  that accounts both for the local  
224 homogeneity and the local distance to the style image, for some images, it is difficult to find a good  
225 compromise between the two constraints. More specifically, we observed that both terms as well as  
226 their relative importance could vary a lot from a problem to another. In the case where the image  
227 is highly locally contrasted, for a black and white landscape picture for instance, the variance term  
228 was sometimes always largely dominant over the distance one. Therefore, there was no possibility  
229 to improve the fidelity of the mapping without producing a uniformly fine-split quadtree on the  
230 largest part of the final image. The opposite situation could also occur, although less often as the  
231 variance tended to have both a higher value and variation over the blocks. In the case where the  
232 content image and the style image were very dissimilar, for instance if the color tones of the two  
233 images would differ greatly, trying to further constrain the scale of the reproduced details would be  
234 limited by the closeness term, and a thin splitting would occur in roughly homogeneous regions of  
the original image.

235 To deal with those cases, we decided to separate the contributions of the two terms. A split therefore  
236 occurs as long as the minimal size hasn't been reached if one of the two constraints, the local  
237 homogeneity or the closeness to the style patches is violated.

238 The new condition that we implemented was therefore:  
239

$$240 \quad C(p_{x_i}, p_{y_i}) = (((\sigma_i > \omega_{var}) \text{ or } (d(p_{x_i}, p_{y_i})) > \omega_{dist}) \text{ and } \tau_i > \gamma_{min}) \text{ or } (\tau_i < \gamma_{max})$$

241 Which can also be written  
242

$$243 \quad C(p_{x_i}, Patches) = (((\sigma_i > \omega_{var}) \text{ or } (d(p_{x_i}, Patches)) > \omega_{dist}) \text{ and } \tau_i > \gamma_{min}) \text{ or } (\tau_i < \gamma_{max})$$

244 Where  $(d(p_{x_i}, Patches))$  is the distance from the the block  $p_{x_i}$  to the closest item from  $Patches$   
245 which represents all the extracted patches from the style image that have the same size as  $p_{x_i}$ .  
246

247 This rule allows us to control more precisely the splitting for a given image by balancing the two  
248 constraints independently to fit our specific needs.  
249

250 On a specific example with large patch variance variation as in 1 we can see that for a same sum  
251 of thresholds, we can finetune the splitting for a given image. In this case, we purposefully chose a  
252 style image that was 'distant' from the content image, so that both thresholds are regularly activated  
253 in the thresholding.

254 Here we can see that by decreasing the variance threshold without modifying the sum of thresholds  
255 we can control the splitting around the content image details. By lowering the variance threshold to  
256 10, we manage to produce splits around the mouth and the outline of the face, and therefore increase  
257 our chances of making the face recognizable in the final image. If we try produce the same effect by  
258 controlling the sum of the thresholds only, we can reduce this global threshold enough to produce  
259 the same details, but in this case we also produce additional splits in the region of the hair, where  
260 almost all splits occur until they reach the minimal splitting size, creating a large uniform zone of  
261 fine-grained splitting which letter often fails to faithfully reproduce the style.  
262

### 263   3.2 Implementation details 264

265 The implemented algorithm for the splitting is close to the the original one, although the condition  
266 for the split is different.  
267

268  $regionSplit(image, \gamma)$  splits the original image in squares of size  $\gamma$  while  $getPatches(image, \gamma)$   
269 retrieves a set of patches of size gamma from the given image.

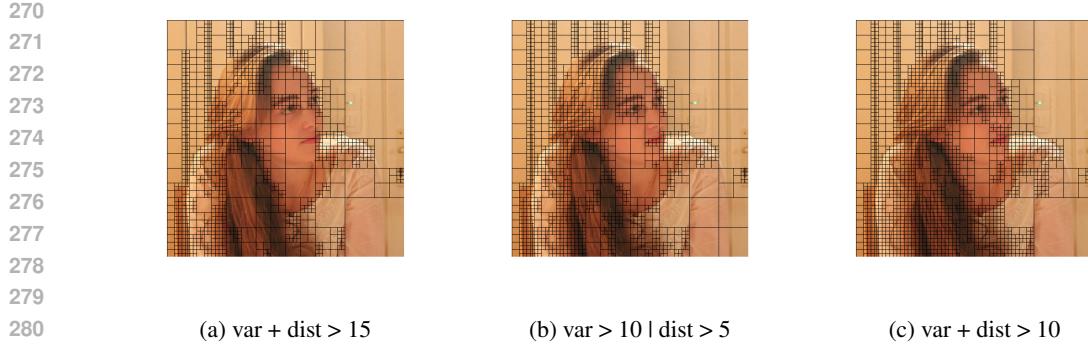


Figure 1: Quadtree splits for various thresholds

---

**Algorithm 1** Split and match algorithm

---

```

285   inputs Images : imcontent, imstyle; Thresholds :  $\omega_{var}, \omega_{dist}$ ; Limitsizes :  $\gamma_{min}, \gamma_{max}$ 
286   outputs Set of blocks :  $R = \{R_i\}_{i=1,n}^n$ ; Assigned labels :  $L = \{L_i\}_{i=1,n}^n$ 
287   Initialization :  $R = regionSplit(imcontent, gamma_{max})$ 
288    $\gamma = \gamma_{max}$ 
289   while  $\gamma \geq \gamma_{min}$  do
290     patches = getPatches(imstyle,  $\gamma$ )
291     for  $R_i \in R$  do where  $size(R_i) = \gamma$ 
292        $\sigma_i = \sqrt{var(R_i)}$ 
293       if  $C(R_i, patches)$ 
294         Split  $R_i$  into four
295          $R = \{R \setminus R_i\} \cup \{R_i^1, R_i^2, R_i^3, R_i^4\}$ 
296       else if
297          $L_i = p_{[(argmin_j(dist(R_i, patches[j])))]}^{imstyle}$ 
298       end if
299     end for
300      $\gamma = \gamma/2$ 
301   end while
302   return
303

```

---

This algorithm outputs the final quad-tree splitting of the original image ( $R$ ), and the labels that map each block to a patch from the style image. The algorithm therefore performs the splitting and mapping operations that strive at operating style transfer. We then have to map the selected style patches to the corresponding region position in the image to obtain the final image.

Several optimizations were performed in order to reduce the computing time, for instance, the getPatches function is only called in the case where one of the blocks at the given size  $\gamma$  passes the threshold test. The distance to the current block is also only computed if the variance threshold is passed, as the variance computation is relatively fast in comparison to the distance computation when the number of patches is large. This is furthermore justified in practice, as the variance threshold is often triggered.

Another important aspect of the algorithm is the sampling of patches from the style image. The initial style image is of a given size, and provides wide possibilities for various patch extraction methods. The first one that we implemented was a simple cut without overlap at each size  $\gamma$  : the style image was cut into contiguous square patches. To augment the number of candidates, we also tried to sample with some overlap between the patches, with a fixed overlap size of  $decal$ . In this case, at a given patch size  $\gamma$  the number of sampled patches for a square style image of size  $imageSize$  is  $\text{floor}(\frac{imageSize - \gamma}{decal})^2$ , therefore, the number of patches increases as  $1/decal^2$ . As the k-nn algorithm has to compute the distance to every single patch, a small  $decal$  negatively impacted the performances of the algorithm. On the other hand, when  $decal$  is above the current patch size  $\gamma$ , less patches are sampled than if they were contiguously sampled, which can account for a speed-up of the k-nn search.

324 Another option would have been to adapt the size of the overlap to  $\gamma$  as for instance to produce a  
325 constant number of patches at each size.  
326

### 327 3.3 Boundary cut 328

329 Experiments on texture synthesis showed that bilinear blending works well if the boundaries are very  
330 similar to one another, but tend to perform poorly if the color or illumination variations are more  
331 drastic. One of the improvements of Efros and Freeman which allowed them to produce visually  
332 coherent textures was using boundary cuts as depicted in [3]. A boundary cut is performed by  
333 computing a surface error that measures the distance at each overlapping pixel. Once this surface is  
334 computed, the cut is operated smoothly by minimizing the global cut price, which is the sum of the  
335 distances for the pixels over which the cut path is drawn. More specifically, in my implementation,  
336 we decided (arbitrarily) to cut to the right of the minimal path, as we cannot cut a pixel in the middle,  
337 and must therefore choose whether to cut at the right or at the left of the minimal path. We could  
338 also have operated a mean on the pixels that correspond the chosen path.  
339

340 To compute the optimal path, we compute the cumulative errors starting from the last row by adding  
341 the current error to the minimal cumulative errors from the previous rows among the three closest  
342 cumulative error candidates.

---

#### 343 Algorithm 2 Compute cumulative errors

---

```
344 inputs Overlaps : overlap1, overlap2
345 outputs Optimal cut path : path
346 errorSurface = getErrorSurface(overlap1, overlap2)
347 cumulativeError(lastRowIndex) = errorSurface(lastRowIndex)
348 for rowIndex from lastRowIndex - 1 to 1 with a step of -1 do
349     for overlapIndex from 1 to 1 overlapSize do
350         cumulativeError(rowIndex, overlapIndex) = min(cumulativeError(rowIndex + 1, overlapIndex - 1),
351             cumulativeError(rowIndex + 1, overlapIndex), cumulativeError(rowIndex + 1, overlapIndex + 1)) + errorSurface(rowIndex, overlapIndex)
352     end for
353 end for
354 return
```

---

355  
356 *getErrorSurface(overlap1, overlap2)* computes the error surface as the sum of square errors  
357 over the three color channels between the two overlaps.

358 Once we have the cumulative errors, we can look at the minimal value on the first row, and find a  
359 way back through the minimal cumulative error path.

360 For the cutting algorithm, the main parameter is the size of the overlapping area, which for convenience  
361 was taken as a multiple of two, such as to allow for the extraction of patches from the style  
362 image with a padding of half the size of the overlap in all directions. In practice, a total width of  
363 overlap of 4 or 6 provided good results in most cases. The compromise that was considered is the  
364 quality of the blending versus the ability to preserve fine-grained details. Increasing the size of the  
365 overlapping area allowed for more flexibility in the cut, which provided more potential cut paths, and  
366 therefore improved the blending. The downside was that those cuts might discard relevant details of  
367 the content which are in the same order of size as the overlap size.

368 It is interesting to notice that, having implemented a boundary cut following as presented by Efros  
369 and Freeman in [3], our method is a generalization of their patch-quilting algorithm for square  
370 images. Their square quilting algorithm is almost the same as ours in the specific case where  
371  $\gamma_{min} = \gamma_{max}$ , and therefore the local variance and patch distance are not taken into account.

### 373 3.4 Regularization 374

375 Running the algorithm without any regularization, we can observe that small repetitive patterns  
376 appear, affecting the global quality of the transfer. We observed that this effect mainly comes from  
377 the same patch of the style image being mapped to neighbors of the content image, sometimes  
repeatedly over an entire region.

To avoid this phenomena, the original paper produced several candidates for each block and produced a result minimizing a likelihood that took into account regularization parameters. In our case, we operate the regularization in a different way, at the same time as we operate the splitting and quilting of the final image.

The first consideration we made was that the limitation of patch repetition could be operated for each quadtree dimension separately, as two patches of different sizes are by definition different. The second consideration was that the main visual defect was the exact repetition of a unique patch. We therefore adapted the distance term to take into account the distance between identic patches.

To achieve this, we modified the distance in the following way

$$d(p_{x_i}, p_y) = \frac{\|p_{x_i} - p_y\|}{\tau_i^2} + \frac{\lambda}{dist(x_i, x_j)/\tau_i}$$

Where d is the euclidian distance and  $x_j$  is the center of the closest block that has been assigned the label that corresponds to the patch  $p_y$  of the style image. In the case where no block has been assigned the candidate label before, the distance of the candidate to the same label in the content image is infinite, and therefore the regularization term equals zero. In this case, the total distance is the same as it was previously defined. Also, setting  $\lambda$  to zero is equivalent to using the original distance.

Taking the inverse of the distance allows for large penalization of close duplicates of labels and low penalization of distant ones.

As measuring the distance for all possible candidates from the style image would be computationally prohibitive, we restricted the computation to a limited set of closest neighbors computed by the k-nn algorithm based purely on the normalized patch distance. This also allowed to avoid the case where a locally irrelevant patch would be selected, because all relevant patches have already been mapped closely. In this case, repetition would probably harm the final result less than an irrelevant patch selection.

In practice, we observed that setting  $\lambda$  to 4 and choosing between k=5 closest neighbors produced good results : it avoided strict repetition while having a limited effect on the quality of the mapping as far as the proximity to the original data was concerned.

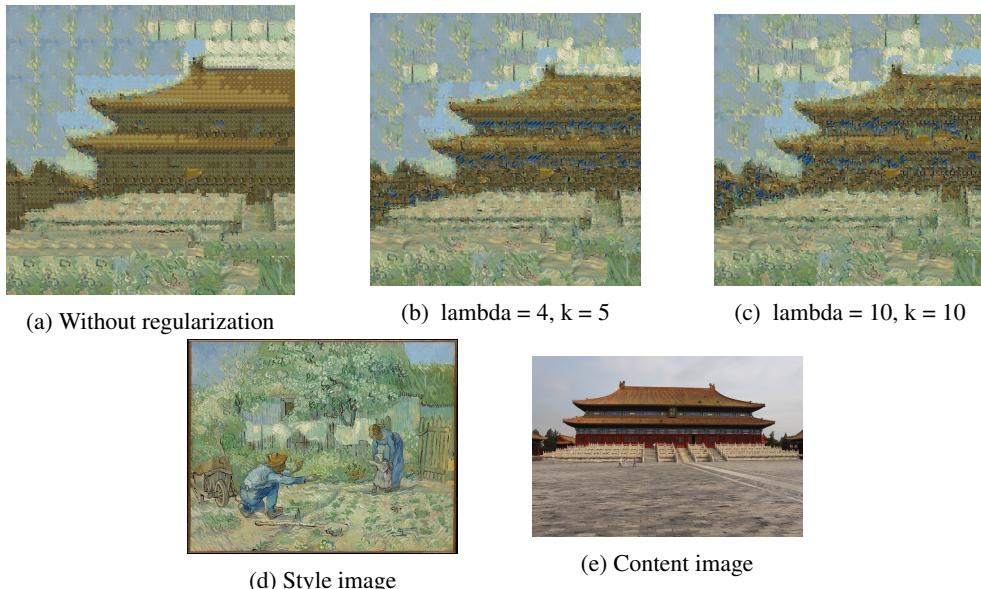


Figure 2: Effect of regularization

We can see looking at 2 an example of final images obtained for a content-style pair for various regularization factors and neighbor numbers. We can see that adding distance regularization sig-

432 significantly reduces the visual patch repetition. Between 2b and 2c there is an overall diminution of  
433 patch repetition. Nevertheless, in the left top part of the sky, we can see the same pattern repeated  
434 several times. Since the regularization is operated on the distance only to the exact same patch, and  
435 sometimes, there are very similar patches among the candidates. In our current implementation,  
436 similar patches are note penalized.

437 The observed limitation of the regularization becomes stronger when we increase *decal* and therefore  
438 sample a larger number of partially overlapping patches. To take into account the risk of having  
439 very similar patches with different labels that might create a visual repetition of patch pattern, in-  
440 stead of taking into account the label of the patches, we could penalize the candidates by taking into  
441 account the similarity to their already quilted neighbors, using normalized pixel distance over the  
442 style patches for instance.

### 444 3.5 Results

445 In 3, we present some examples of style transfer operated from pictures of paintings that are freely  
446 available for download and use on the website of the Museum of Modern Arts. We chose *View*  
447 from the artist's window by Robert Blum, *German Landscape with View towards a Broad Valley* by  
448 Fritz Petzholdt and *Study for Portrait of Mrs. Anna E. Little* by Childe Hassam for the variety of  
449 the represented styles. The content pictures were also taken to be different from one another. We  
450 therefore chose a photo of architectural details of houses, a landscape, and a close-up on a flower.  
451 Those results are presented next to some results that were obtained by running the style transfer  
452 through a freely available implementation of the neural style transfer algorithm [6] that produce  
453 similar results in quality to [5] but with a significant speed-up [7]. The neural style algorithm  
454 starts from a random noise image and operates several iterations to minimize the loss that accounts  
455 for style consistency and content preservation. In the presented examples, 1000 iterations over  
456 the neural network were performed for each image, and in all cases, no significant improvement  
457 was observed during the last 300 iterations, which indicates the convergence of the neural transfer  
458 algorithm.

459 For our method, the patch dimensions were set to  $\gamma_{min} = 8$ ,  $\gamma_{max} = 64$ , the thresholds were set to  
460  $\omega_{var} = 18$   $\omega_{dist} = 5$ . The regularization coefficient was set to  $\lambda = 4$ , and the *decal* that regulates  
461 the size of the candidate patch set was set to 32. For the cut, each patch was sampled with a padding  
462 of 3 pixels, producing a width for overlap region of 6 pixels.

463  
464  
465  
466  
467  
468  
469  
470  
471  
472  
473  
474  
475  
476  
477  
478  
479  
480  
481  
482  
483  
484  
485

486  
487  
488  
489  
490  
491  
492  
493  
494  
495  
496  
497

### Style images



(a)



(b)



(c)

498  
499  
500  
501  
502  
503  
504  
505  
506  
507  
508  
509

### Content images



(d)



(e)



(f)

510  
511  
512  
513  
514  
515  
516  
517  
518  
519  
520  
521

### Our results



(g)



(h)



(i)

522  
523  
524  
525  
526  
527  
528  
529  
530  
531  
532  
533  
534

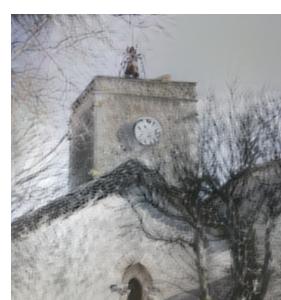
### Neural style results



(j)



(k)



(l)

535  
536  
537  
538  
539

Figure 3: Comparison of style transfer results

We observe that both methods fulfill to some extent the style transfer task and that each method has its own strengths and weaknesses. We can see that for our method, the cuts are still sometimes

540 visible, for instance in the case of the flower 3e. And the patch repetition phenomena is slightly  
 541 visible in 3i. In 3g, we can see an example where content of the style image has been transferred  
 542 with the style, as we can see that the fence from the original style image 3a is visible in the sky. For  
 543 the neural style transfer, we can observe a slight deformation in 3l that is not linked to a style-related  
 544 deformation, and in 3j we observe that the painting texture was poorly transferred on the sky, leaving  
 545 it very similar to the original image, although the rest of the image was successfully transformed.  
 546

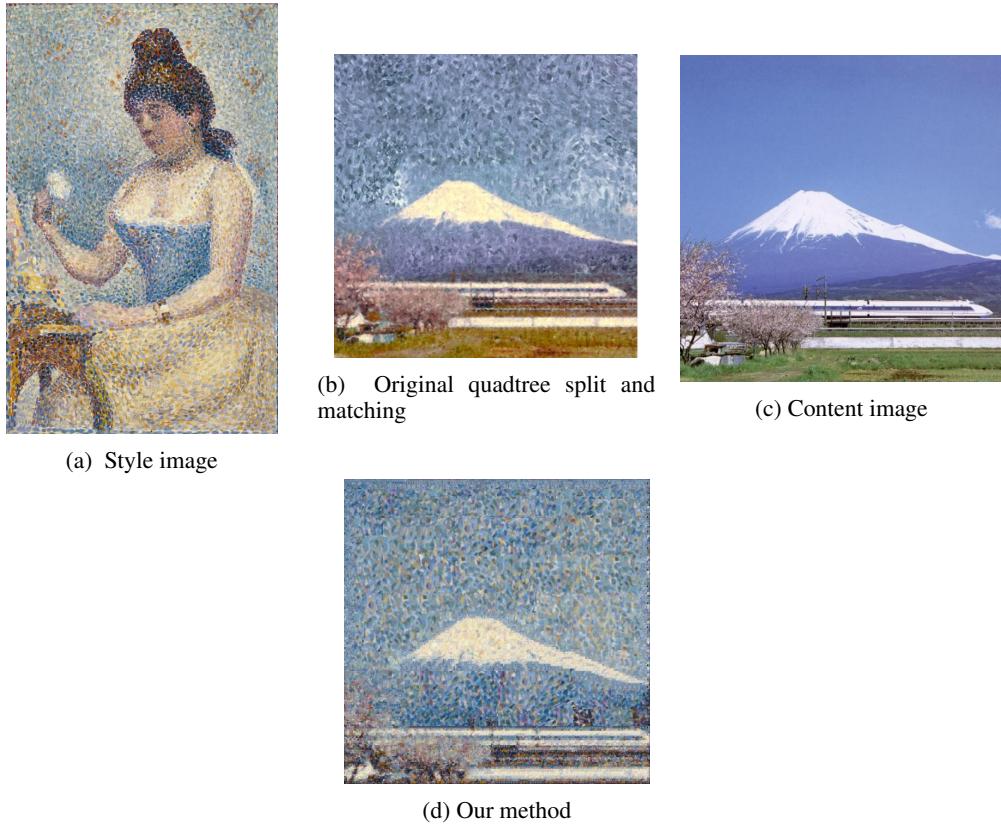


Figure 4: Comparison with original method

575 We can also see, looking at 4 that our implementation performs comparatively to the original im-  
 576 plementation. Although in our case, despite the flexibility allowed by irregular patch cutting, seams  
 577 are sometimes still visible. This observation justifies the presence in the original implementation of  
 578 a regularization term that takes into account the overlap distance. In the original method, it seems  
 579 that around the limits of the mountain, slightly excessive splitting occurred, producing an almost  
 580 homogeneous region. In our case, by appropriately fixing a variance threshold of 15 and a distance  
 581 threshold of 5, we manage to avoid this effect. As the variance around the region is very low, this  
 582 effect on the original image was probably triggered by a larger distance for light blue patches from  
 583 the content image to the style image. Splitting the two terms allows us to reduce the splitting around  
 584 this area, without penalizing the sharpness of the splitting around the edges of the mountain.

### 585 3.6 Limitations

587 While experimenting, we have noticed the occasional appearance of artefacts produced by our algo-  
 588 rithm.

#### 589 Color artefacts

591 Because of the choice of the norm, color artefacts can appear. As the minimized distance that  
 592 determines the closest neighbor is the sum of square over each separate channel, no color continuity  
 593 is insured at the border between two patches. This problem appears specifically when the content  
 594 image and style image have different dominant colors. If for instance the content image contains

594 a light yellow block as on the forehead in 5c, and the style image does not contain any patch in  
 595 the given tone, a minor difference in neighbor blocks with minor color variation can be mapped  
 596 to style patches of various colors. This produces a color artefact as in 5 where a block on the  
 597 forehead is mapped to a blue patch while its neighbors are mapped to red patches, creating a visible  
 598 discontinuity.



600  
 601  
 602  
 603  
 604  
 605  
 606  
 607  
 608  
 609  
 610  
 611  
 612  
 613  
 614  
 615  
 616  
 617  
 618  
 619  
 620  
 621  
 622  
 623  
 624  
 625  
 626  
 627  
 628  
 629  
 630  
 631  
 632  
 633  
 634  
 635  
 636  
 637  
 638  
 639  
 640  
 641  
 642  
 643  
 644  
 645  
 646  
 647

(a) Resulting image - without regularization

(b) Style image

(c) Content image

Figure 5: Example of color artefact

### Variance artefact

Artefacts can also appear because of the fact that the thresholds are shared for all scales  $\gamma$ . In 6a we can see that the forehead of the man is cut, although the limit of the forehead is a zone of high variance. But given the fact that the variance is computed over a large patch, the fact that this patch is otherwise very homogeneous rules against a split, which produces this uncanny sharp cut. When high variance is very localized on a patch at a high scale, the detail might nevertheless get lost.

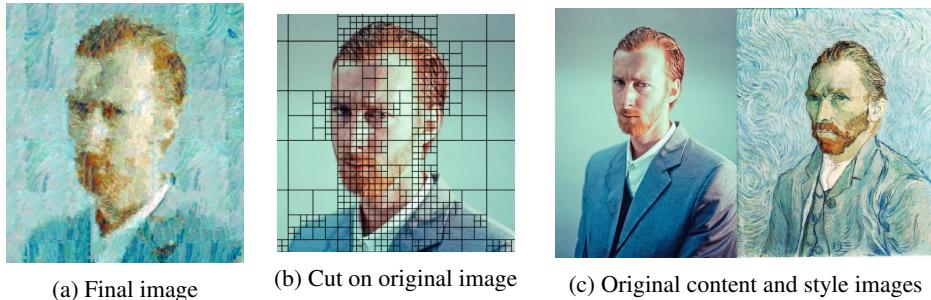


Figure 6: Example of variance artefact

### 3.7 Conclusion

In this work, we have tried to build upon the adaptative quadtree quilting method proposed by Frigo et al [4]. We implemented an algorithm inspired by the original article trying to address its potential shortcomings. In order to achieve better control over the style transfer, we modified the quadtree splitting condition to take into account the local homogeneity of the image and the quality of the potential match separately. Our implementation is therefore closer to standard texture transfer methods that synthesize new pixels by considering the already synthesized ones, rather than taking into account a global constraint over the probability distribution of the samples. We tried to avoid the repetition of similar patches by introducing a regularization directly in the splitting condition in order to progressively build the new image while computing the quadtree splits instead of separating the two steps, we also implementing a boundary cut to diminish the visibility of the seams in the final image. Compared to the original method, the quality of the overlap is not directly considered to choose the best candidate patch. This constraint could be implemented in our method at split-and-quilting time, by adding a second regularization term that favors candidates that have a matching overlap on the already synthesized method. Compared to state of the art results provided by the

648 neural style transfer method, we produce results that have different shortcomings. The neural style  
649 approach is intrinsically free of the seam-blending problem as the image is processed as a whole  
650 rather than decomposed into its parts, but can produce artefacts especially at lower scales by en-  
651 forcing correlations that produce visible wavelets that are not style-related. Emphasizing overlap  
652 similarity could be a solution to solve artefact occurrences in our method. Enforcing overlap simi-  
653 larity by adding a regularization term in the computed distance could solve the variance artefact  
654 by patching patches that have similar borders near to the patches that have been cut despite local  
655 variance if the local variance appears at the border, which is usually the case. It would also play in  
656 favor of color continuity if the regularization term takes into account each channel so as to encourage  
657 similarity between overlapping pixel colors. To improve our algorithm’s performance, we could also  
658 provide patches sampled from various images in the same style, from a given painter for instance.  
659 This would increase the size of the set of potential candidates, without the drawbacks of artificially  
660 increasing the size of the set by sampling various overlapping crops in the same image. In this case,  
661 looking at the selected patches over various transfer problems, we could to some extent define the  
662 style of a painting as the set of candidate patches of various sizes. We could even restrict ourselves  
663 to patches that are selected during the quilting operations, as the selected patches would be close to  
664 low variance patches in the content image, and would therefore likely not contain content details. If  
665 this set of patches consistently allowed us to transfer the style to various content images, we could  
666 consider that the set of patches entirely characterizes the style. Overall, style transfer could further  
667 benefit from insights on what characterizes the style of an image, by analyzing the constraints over  
668 a neural network activations that allow us to capture the style, or by analyzing other algorithms that  
669 aim explicitly at operating style transfer.

## 670 References

- [1] Michael Ashikhmin. Synthesizing natural textures. *ACM Symposium on Interactive 3D Graphics*, pages 217–226, 2001.
- [2] Michael Ashikhmin. Fast texture transfer. *IEEE Computer Graphics and Applications*, pages 38–43, July/August 2003.
- [3] Alexei A. Efros and William T. Freeman. Image quilting for texture synthesis and transfer. *Proceedings of SIGGRAPH 2001*, pages 341–346, August 2001.
- [4] Oriel Frigo, Neus Sabater, Julie Delon, and Pierre Hellier. Split and match: Example-based adaptive patch sampling for unsupervised style transfer. *CVPR*, 2016.
- [5] Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge. A neural algorithm of artistic style. *ArXiv e-prints*, 2015.
- [6] Justin Johnson. neural-style. <https://github.com/jcjohnson/neural-style>, 2015.
- [7] Justin Johnson, Alexandre Alahi, and Li Fei-Fei. Perceptual losses for real-time style transfer and super-resolution. In *European Conference on Computer Vision*, 2016.
- [8] Li-Yi Wei and Marc Levoy. Fast texture synthesis using tree-structured vector quantization. *SIGGRAPH*, page 479–488, 2000.